

Project Title : Customer Sentiment Analysis and Satisfaction Prediction from Support Tickets

language : Machine learning, python, SQL, Excel

Tools : VS code, Jupyter notebook, Google Colab (Collaboratory)

Dataset : Dataset is available in the given link. <https://drive.google.com/file/d/1DRdLKOinSNuoMwVyFGH86f3xEhkZMrz6/view>

Project Overview

The primary objective of this project is to predict customer satisfaction using the historical data contained within the Customer Support Ticket Dataset. This will be achieved by applying machine learning algorithms to analyze the various factors and features present in the ticket information that have an influence on how satisfied a customer is with the support they receive. The ultimate goal is to build a predictive model capable of estimating customer satisfaction levels based on the characteristics of their support interaction.

About Dataset

The Customer Support Ticket Dataset contains records of customer inquiries related to tech products, covering issues like hardware, software bugs, network problems, account access, and data loss. It includes details about the customer (name, email - anonymized domain, age, gender), the product purchased, the purchase date, and specifics of the support ticket (type, subject, description, status, resolution, priority, channel). Crucially, for closed tickets, it also provides the first response time, time to resolution, and the customer satisfaction rating (on a scale of 1 to 5). This dataset is suitable for analyzing support trends, applying Natural Language Processing to ticket text, predicting customer satisfaction and resolution times, segmenting customers, and building recommendation systems for support resources.

Features Description:

Ticket ID: A unique identifier for each individual support ticket.

Customer Name: The name of the customer who submitted the ticket.

Customer Email: The customer's email address (domain anonymized for privacy).

Customer Age: The age of the customer.

Customer Gender: The gender of the customer.

Product Purchased: The specific tech product the customer bought.

Date of Purchase: The date when the customer purchased the product.

Ticket Type: The broad category of the support request (e.g., technical issue).

Ticket Subject: A brief summary or topic of the customer's ticket.

Ticket Description: A more detailed explanation of the customer's problem or question.

Ticket Status: The current state of the ticket (e.g., open, closed).

Resolution: The solution or action taken to close the ticket.

Ticket Priority: The assigned level of urgency for the ticket.

Ticket Channel: How the customer contacted support (e.g., email, chat).

First Response Time: The duration until the initial response was sent to the customer.

Time to Resolution: The total time taken to resolve the ticket.

Customer Satisfaction Rating: The customer's feedback on their satisfaction with the resolution (1-5 scale).

Use Cases of such dataset:

Customer Support Analysis: This allows for data-driven improvements to support operations by understanding what issues are most frequent, how efficiently they are handled, and where processes can be optimized.

Natural Language Processing (NLP): Leveraging the text data (subject and description) can lead to automation and deeper understanding of customer needs and sentiments at scale.

Customer Satisfaction Prediction: Proactively identifying customers likely to be dissatisfied enables timely interventions and can improve overall customer loyalty.

Ticket Resolution Time Prediction: Accurate predictions can help manage customer expectations, allocate resources effectively, and potentially identify tickets that might require escalation.

Customer Segmentation: Understanding different customer groups based on their support interactions allows for tailored support strategies and potentially personalized product or service offerings.

Recommender Systems: Providing relevant solutions or product suggestions within the support process can improve efficiency and customer self-service capabilities.

Data Preprocessing:

This crucial step involves cleaning and preparing the raw data for modeling.

The sample code demonstrates:

Loading the dataset using pandas.

Displaying basic information about the data (data types, non-null values).

Handling missing values by dropping rows containing any NaN values.

Encoding categorical (object) columns into numerical representations using LabelEncoder.

Exploratory Data Analysis (EDA):

While not explicitly shown in the provided code snippet, EDA is a vital step that typically precedes feature engineering and model building. It involves visualizing data distributions, identifying patterns, and understanding relationships between variables using libraries like matplotlib and seaborn. This helps in gaining insights into the data and informing subsequent steps.

Feature Engineering:

This involves selecting, transforming, and creating new features from the existing data that might improve the performance of the machine learning model. The sample code shows a basic step of defining the feature set (X) by dropping 'CustomerID' and the target variable ('Overall Satisfaction'), and defining the target variable (y). More advanced feature engineering could involve creating interaction terms, polynomial features, or extracting information from date/time columns.

Model Building:

This step involves selecting a suitable machine learning algorithm and training it on the prepared data. The sample code uses a RandomForestClassifier, a popular ensemble method known for its good performance on various classification tasks. The data is split into training and testing sets to evaluate the model's ability to generalize to unseen data.

Model Evaluation:

After training, the model's performance is evaluated on the test set using appropriate metrics. The sample code demonstrates the use of:

accuracy_score:

The proportion of correctly classified instances. `classification_report`: Provides precision, recall, F1-score, and support for each class.

confusion_matrix:

A table that summarizes the model's predictions against the actual values, showing true positives, true negatives, false positives, and false negatives.

Visualization:

This step involves presenting the results and insights in a visual format for better understanding and communication. The sample code shows an example of visualizing feature importance from the trained Random Forest model, highlighting the top 10 features that contributed most to the predictions.

```
from google.colab import files
uploaded = files.upload()
```



Choose Files customer_s...t_tickets.csv

- **customer_support_tickets.csv**(text/csv) - 3945533 bytes, last modified: 5/7/2025 - 100% done
Saving customer_support_tickets.csv to customer_support_tickets (6).csv

```
import pandas as pd
import io
```

```
if uploaded:
    filename = list(uploaded.keys())[0] # Get the first uploaded filename
    print(f"Uploaded filename: {filename}")
    try:
        df = pd.read_csv(io.BytesIO(uploaded[filename]))
        print("\nDataset loaded successfully!")
        print(df.head()) # Display the first few rows
```

```

print(df.info()) # Display dataset information
except Exception as e:
    print(f"\nError loading the dataset: {e}")
else:
    print("\nNo file was uploaded.")

0      Product setup
1  Peripheral compatibility
2      Network problem
3      Account access
4      Data loss

Ticket Description \
0 I'm having an issue with the {product_purchase...
1 I'm having an issue with the {product_purchase...
2 I'm facing a problem with my {product_purchase...
3 I'm having an issue with the {product_purchase...
4 I'm having an issue with the {product_purchase...

Ticket Status      Resolution \
0 Pending Customer Response      NaN
1 Pending Customer Response      NaN
2 Closed Case maybe show recently my computer follow.
3 Closed Try capital clearly never color toward story.
4 Closed West decision evidence bit.

Ticket Priority Ticket Channel First Response Time Time to Resolution \
0 Critical Social media 2023-06-01 12:15:36      NaN
1 Critical Chat 2023-06-01 16:45:38      NaN
2 Low Social media 2023-06-01 11:14:38 2023-06-01 18:05:38
3 Low Social media 2023-06-01 07:29:40 2023-06-01 01:57:40
4 Low Email 2023-06-01 00:12:42 2023-06-01 19:53:42

Customer Satisfaction Rating
0      NaN
1      NaN
2      3.0
3      3.0
4      1.0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8469 entries, 0 to 8468
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Ticket ID              8469 non-null  int64
1   Customer Name          8469 non-null  object
2   Customer Email         8469 non-null  object
3   Customer Age           8469 non-null  int64
4   Customer Gender        8469 non-null  object
5   Product Purchased      8469 non-null  object
6   Date of Purchase       8469 non-null  object
7   Ticket Type            8469 non-null  object
8   Ticket Subject         8469 non-null  object
9   Ticket Description     8469 non-null  object
10  Ticket Status          8469 non-null  object
11  Resolution             2769 non-null  object
12  Ticket Priority        8469 non-null  object
13  Ticket Channel        8469 non-null  object
14  First Response Time    5650 non-null  object
15  Time to Resolution     2769 non-null  object
16  Customer Satisfaction Rating 2769 non-null  float64
dtypes: float64(1), int64(2), object(14)
memory usage: 1.1+ MB
None

```

```

# Assuming your preprocessed DataFrame is named 'df_cleaned'
print(df_cleaned.columns)

```

```

Index(['Ticket ID', 'Customer Name', 'Customer Email', 'Customer Age',
      'Customer Gender', 'Product Purchased', 'Date of Purchase',
      'Ticket Type', 'Ticket Subject', 'Ticket Description', 'Ticket Status',
      'Resolution', 'Ticket Priority', 'Ticket Channel',
      'First Response Time', 'Time to Resolution',
      'Customer Satisfaction Rating'],
      dtype='object')

```

```


# Analyze customer support ticket trends
# Identify common issues
common_issues = df['Ticket Subject'].value_counts().head(10)
print("Top 10 Common Issues:")
print(common_issues)

```

```
# Plotting ticket trends over time
df['Date of Purchase'] = pd.to_datetime(df['Date of Purchase'])
df['YearMonth'] = df['Date of Purchase'].dt.to_period('M')
ticket_trends = df.groupby('YearMonth').size()
plt.figure(figsize=(10, 6))
ticket_trends.plot(kind='line', marker='o')
plt.title('Customer Support Ticket Trends Over Time')
plt.xlabel('Year-Month')
plt.ylabel('Number of Tickets')
plt.grid(True)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

↗ Top 10 Common Issues:

Ticket Subject	
Refund request	576
Software bug	574
Product compatibility	567
Delivery problem	561
Hardware issue	547
Battery life	542
Network problem	539
Installation support	530
Product setup	529
Payment issue	526
Name: count, dtype: int64	



```
# Assuming your loaded DataFrame is named 'df'
```

```
# Define the target variable (Customer Satisfaction Rating)
y = df['Customer Satisfaction Rating']
```

```
# Define the features (X) - select the columns you want to use for prediction
# You'll need to choose these based on your understanding of the data and EDA.
```

```
# Here's an example using some potentially relevant columns:
```

```
features = ['Customer Age', 'Product Purchased', 'Ticket Type', 'Ticket Priority', 'Ticket Channel', 'Ticket Subject', 'Ticket Description']
X = df[features]
```

```
# If you've already done some preprocessing and your DataFrame is 'df_cleaned', use that instead:
```

```
# y = df_cleaned['Customer Satisfaction Rating']
```

```
# features = [...]
```

```
# X = df_cleaned[features]
```

```
print("Features (X) shape:", X.shape)
```

```
print("Target (y) shape:", y.shape)
```

```

Features (X) shape: (8469, 7)
Target (y) shape: (8469,)

```

```
# Assuming your loaded DataFrame is named 'df'
```

```
# Segment based on ticket types
ticket_type_segmentation = df.groupby('Ticket Type').size()
print("\nSegmentation based on Ticket Types:")
print(ticket_type_segmentation)
```

```
# Segment based on satisfaction levels
satisfaction_segmentation = df.groupby('Customer Satisfaction Rating').size()
print("\nSegmentation based on Customer Satisfaction Levels:")
print(satisfaction_segmentation)
```

```

Segmentation based on Ticket Types:
Ticket Type
Billing inquiry      1634
Cancellation request 1695
Product inquiry      1641
Refund request       1752
Technical issue      1747
dtype: int64

```

```

Segmentation based on Customer Satisfaction Levels:
Customer Satisfaction Rating
1.0    553
2.0    549
3.0    580
4.0    543
5.0    544
dtype: int64

```

```

import seaborn as sns
import matplotlib.pyplot as plt

```


```
# Assuming your loaded DataFrame is named 'df'
```

```
# Set up the plotting aesthetics
sns.set(style="whitegrid")
```

```

# Customer Satisfaction Distribution
plt.figure(figsize=(10, 6))
sns.histplot(df['Customer Satisfaction Rating'].dropna(), bins=5, kde=True, color='skyblue')
plt.title('Customer Satisfaction Distribution')
plt.xlabel('Satisfaction Rating')
plt.ylabel('Frequency')
plt.show()


```



```
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming your loaded DataFrame is named 'df'


# Ticket Status Distribution
ticket_status_distribution = df['Ticket Status'].value_counts()
plt.figure(figsize=(8, 8))
plt.pie(ticket_status_distribution,
        labels=ticket_status_distribution.index,
        autopct='%1.1f%%',
        colors=sns.color_palette('pastel'),
        startangle=140)
plt.title('Ticket Status Distribution')
plt.axis('equal')
plt.show()
```



```
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming your loaded DataFrame is named 'df'

# Customer Age Distribution
plt.figure(figsize=(10, 6))
sns.histplot(df['Customer Age'].dropna(), bins=20, kde=True, color='salmon')
plt.title('Customer Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```




```
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming your loaded DataFrame is named 'df'

# Customer Gender Distribution
customer_gender_distribution = df['Customer Gender'].value_counts()
plt.figure(figsize=(8, 8))
plt.pie(customer_gender_distribution,
        labels=customer_gender_distribution.index,
        autopct='%1.1f%%',
        colors=sns.color_palette('Set2'),
        startangle=90)
plt.title('Customer Gender Distribution')
plt.axis('equal')
plt.show()
```





Customer Gender Distribution



```
import matplotlib.pyplot as plt
import seaborn as sns

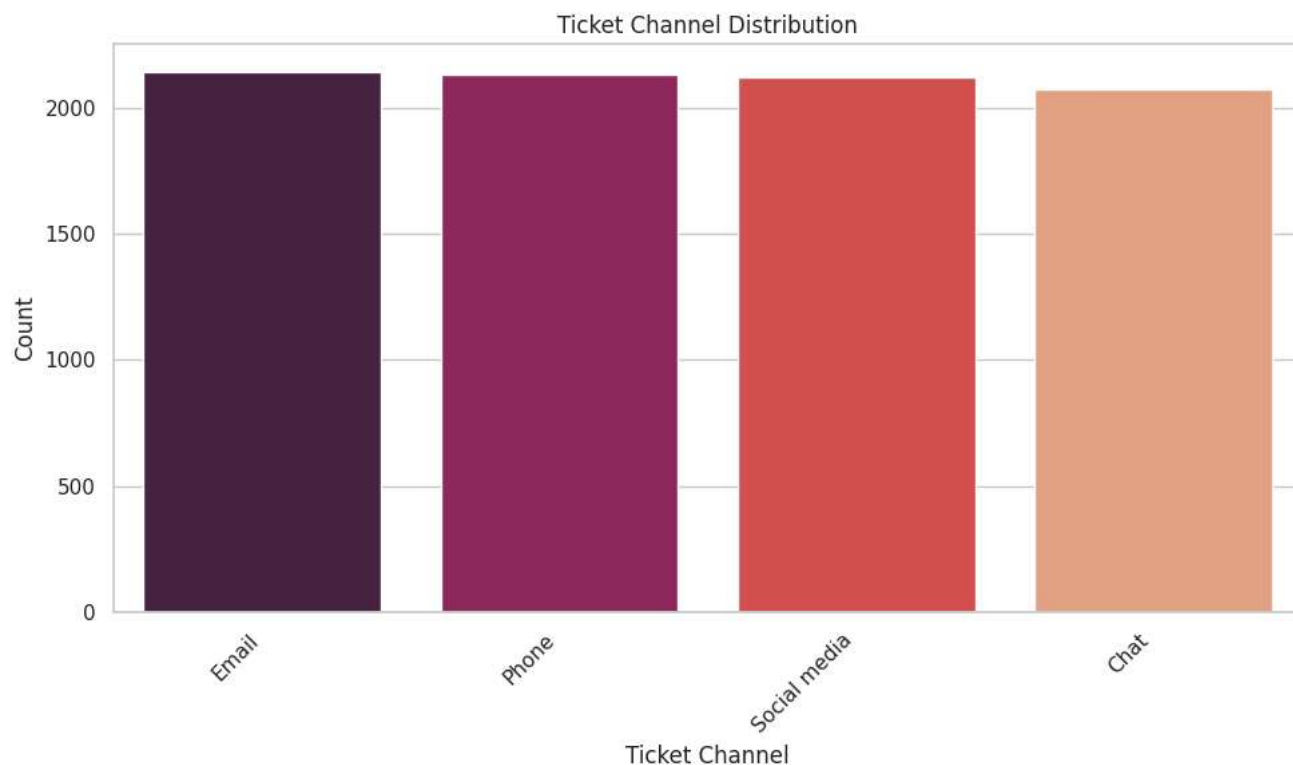
# Assuming your loaded DataFrame is named 'df'

# Ticket Channel Distribution
ticket_channel_distribution = df['Ticket Channel'].value_counts()
plt.figure(figsize=(10, 6))
sns.barplot(x=ticket_channel_distribution.index,
            y=ticket_channel_distribution,
            palette='rocket')
plt.title('Ticket Channel Distribution')
plt.xlabel('Ticket Channel')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Added ha='right' for better label alignment
plt.tight_layout() # Added to prevent labels from overlapping
plt.show()
```

 <ipython-input-37-bc305f982300>:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.barplot(x=ticket_channel_distribution.index,
```



```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Assuming your loaded DataFrame is named 'df'
```

```
# Chart 1: Average Customer Satisfaction by Gender (Bar Plot)
```

```
average_satisfaction = df.groupby('Customer Gender')['Customer Satisfaction Rating'].mean().reset_index()
```

```
plt.figure(figsize=(8, 6))
```

```
sns.barplot(x='Customer Gender', y='Customer Satisfaction Rating', data=average_satisfaction, palette='muted', order=['Male', 'Female', 'Oth
```


```
plt.title('Average Customer Satisfaction by Gender')
```

```
plt.xlabel('Gender')
```

```
plt.ylabel('Average Satisfaction Rating')
```


```
plt.ylim(1, 5) # Adjust y-axis limit if needed
```

```
plt.show()
```

 <ipython-input-38-8fa93097f6b1>:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.barplot(x='Customer Gender', y='Customer Satisfaction Rating', data=average_satisfaction, palette='muted', order=['Male', 'Female'])
```



```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Assuming your loaded DataFrame is named 'df'
```

```
# Product Purchased Distribution
```

```
plt.figure(figsize=(10, 6))
```

```
product_purchased_distribution = df['Product Purchased'].value_counts().head(10)
```

```
sns.barplot(y=product_purchased_distribution.index,
            x=product_purchased_distribution,
            palette='magma')
```


```
plt.title('Top 10 Products Purchased')
```

```
plt.xlabel('Count')
```

```
plt.ylabel('Product')
```

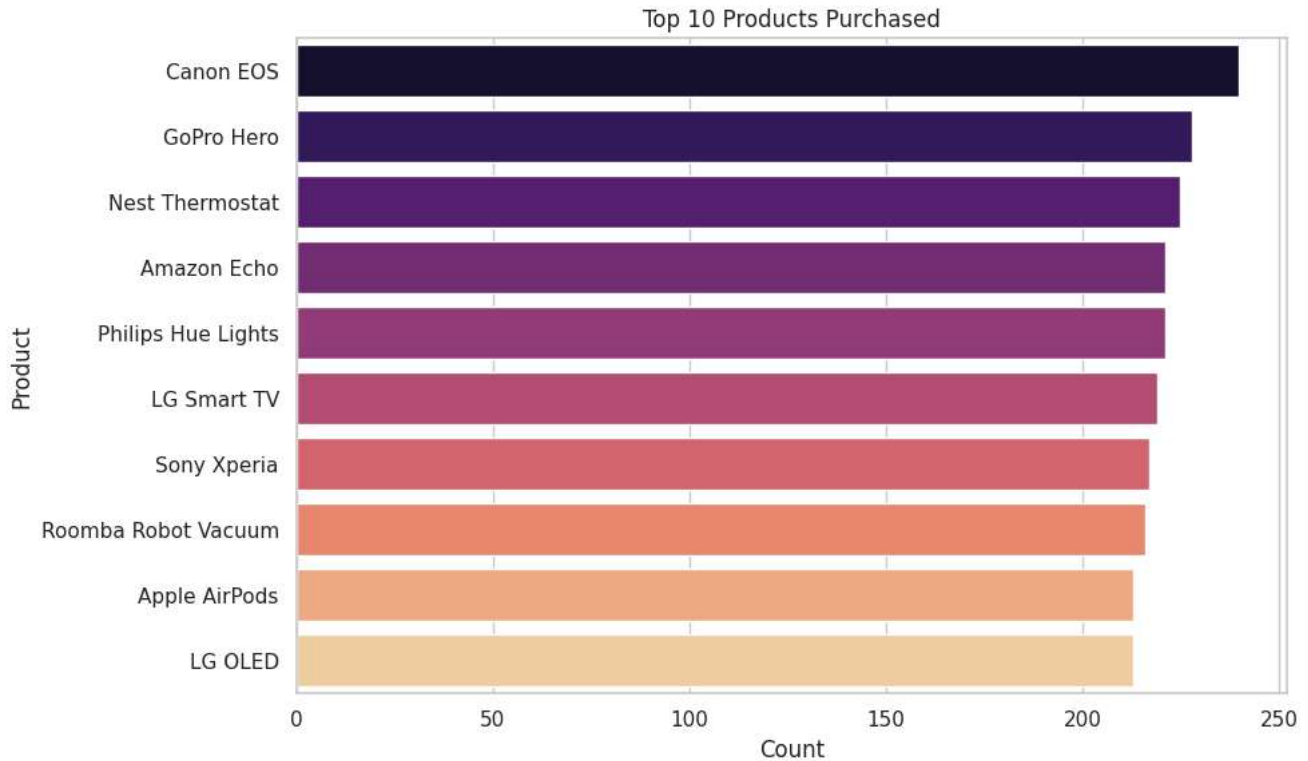
```
plt.tight_layout() # Added to prevent labels from overlapping
```

```
plt.show()
```

 <ipython-input-39-134083decd9a>:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend`

```
sns.barplot(y=product_purchased_distribution.index,
```



```
import matplotlib.pyplot as plt
```

```
# Assuming your loaded DataFrame is named 'df'
```

```
# Chart 2: Top Items Purchased by Gender (Horizontal Bar Chart)
```

```
plt.figure(figsize=(15, 6))
```

```
# Top Items Purchased by Males
```

```
plt.subplot(1, 3, 1)
```

```
top_items_male = df[df['Customer Gender'] == 'Male']['Product Purchased'].value_counts().head(5)
```

```
top_items_male.plot(kind='barh', color='skyblue')
```

```
plt.title('Top Items Purchased by Males')
```

```
plt.xlabel('Count')
```

```
plt.ylabel('Product')
```

```
# Top Items Purchased by Females
```

```
plt.subplot(1, 3, 2)
```

```
top_items_female = df[df['Customer Gender'] == 'Female']['Product Purchased'].value_counts().head(5)
```

```
top_items_female.plot(kind='barh', color='salmon')
```

```
plt.title('Top Items Purchased by Females')
```

```
plt.xlabel('Count')
```

```
plt.ylabel('Product')
```

```
# Top Items Purchased by Other Gender
```

```
plt.subplot(1, 3, 3)
```

```
top_items_other = df[df['Customer Gender'] == 'Other']['Product Purchased'].value_counts().head(5)
```

```
top_items_other.plot(kind='barh', color='lightgreen')
```

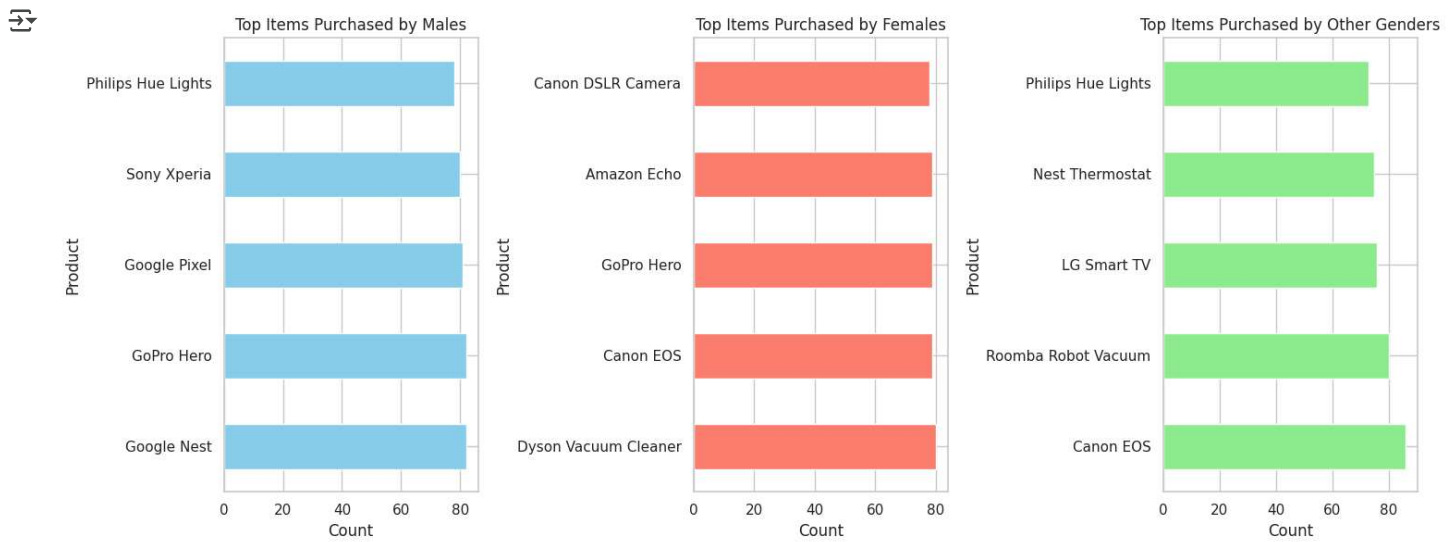
```
plt.title('Top Items Purchased by Other Genders')
```

```
plt.xlabel('Count')
```

```
plt.ylabel('Product')
```

```
plt.tight_layout()
```

```
plt.show()
```

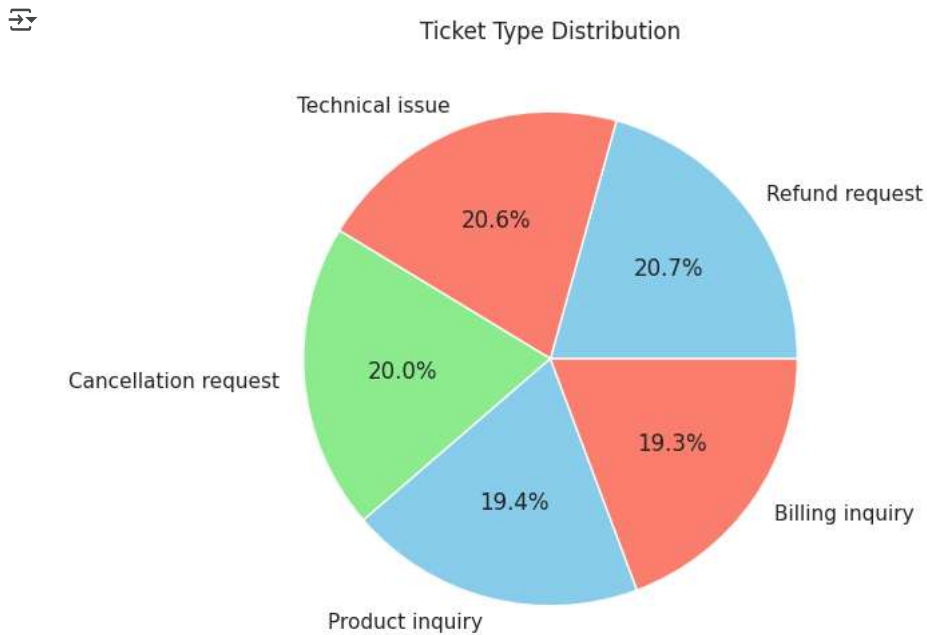


```
import matplotlib.pyplot as plt

# Assuming your loaded DataFrame is named 'df'

# Count ticket types
ticket_type_distribution = df['Ticket Type'].value_counts()

# Plot
plt.figure(figsize=(8, 6))
ticket_type_distribution.plot(kind='pie', autopct='%1.1f%%',
                             colors=['skyblue', 'salmon', 'lightgreen'])
plt.title('Ticket Type Distribution')
plt.ylabel('') # Remove the default 'Ticket Type' label on the y-axis
plt.show()
```



```
import matplotlib.pyplot as plt

# Assuming your loaded DataFrame is named 'df'

# Count ticket priorities
priority_distribution = df['Ticket Priority'].value_counts()
```

```
# Plot
plt.figure(figsize=(8, 6))
priority_distribution.plot(kind='pie', autopct='%1.1f%%',
                           colors=['lightblue', 'lightgreen', 'lightsalmon', 'skyblue'])
```